

Relying on the Unreliable: The Impact of Language Models’ Reluctance to Express Uncertainty

Kaitlyn Zhou^{1,4} Jena D. Hwang⁴ Xiang Ren^{2,4} Maarten Sap^{3,4}

¹Stanford University, ²University of Southern California,

³Carnegie Mellon University, ⁴Allen Institute for AI

katezhou@stanford.edu

Abstract

As natural language becomes the default interface for human-AI interaction, there is a need for LMs to appropriately communicate uncertainties in downstream applications. In this work, we investigate how LMs incorporate confidence in responses via natural language and how downstream users behave in response to LM-articulated uncertainties. We examine publicly deployed models and find that LMs are reluctant to express uncertainties when answering questions even when they produce incorrect responses. LMs can be explicitly prompted to express confidences, but tend to be overconfident, resulting in high error rates (an average of 47%) among confident responses. We test the risks of LM overconfidence by conducting human experiments and show that users rely heavily on LM generations, whether or not they are marked by certainty. Lastly, we investigate the preference-annotated datasets used in post training alignment and find that humans are biased against texts with *uncertainty*. Our work highlights new safety harms facing human-LM interactions and proposes design recommendations and mitigating strategies moving forward.

1 Introduction

Natural language is becoming the default interface for humans to engage with artificial intelligence systems whether it be information seeking, summarization, or image captioning (Bommasani et al., 2021; Brown et al., 2020; Ouyang et al., 2022). As input, natural language serves as a rich and versatile medium, enabling users to articulate intricate tasks and inquiries effectively. As for output, natural language provides an opportunity for language models (LMs) to generate not only informative, but nuanced responses that better support the collaboration between humans and AI.

A pivotal aspect of fostering reliable human-AI interactions lies in the apt communication of model

Question: What is the capital of Mauritania?		Answer: Nouakchott
LM Expressions of Confidence		Human Interpretations
Plain Statement	∅ It's Nouakchott.	⚡⚡⚡⚡⚡⚡⚡⚡⚡⚡
Strengtheners	I'm 100% certain it's Nouakchott.	⚡⚡⚡⚡⚡⚡⚡⚡⚡⚡
Weakeners	I'm not sure, maybe it's Nouakchott.	⚡⚡⚡⚡⚡⚡⚡⚡⚡⚡

⚡ Rely on LM ⚡ Rely on Self

Figure 1: Overview of experiments on human interpretations of epistemic markers. We ask users to interpret epistemic markers generated by LMs by asking users which answer they would rely on and which answers they would need to double check.

uncertainties (Cai et al., 2019; Kizilcec, 2016; DeArteaga et al., 2020), typically defined as the probability assigned to a model’s prediction. Recent work in language generation reflects a shift towards using natural language as a means to convey model confidences (e.g., “I’m fairly confident it’s”, “According to Wikipedia it’s” Mielke et al., 2020; Lin et al., 2022; Zhou et al., 2023). Such features are known in the linguistics literature as *epistemic markers*, which serve to convey a speaker’s stance and commitment, thus supporting human communication and decision making (Babrow et al., 1998; Brashers et al., 2000; Tseng and Zhang, 2023). Since epistemic markers play an important role in human-human communication and decision making (Budescu et al., 1988; Windschitl and Wells, 1996; Druzdzel, 1989), we hypothesize that the use of these markers by LMs will also have an impact on human-AI interactions.

Our work begins with an examination of how LMs communicate uncertainties to end-users in realistic information seeking scenarios (§3). Specifically, we elicit responses from popular, publicly-deployed models including GPT, LLaMA-2, and Claude by prompting them to provide epistemic markers when answering multiple choice questions (Brown et al., 2020; OpenAI, 2023; Touvron et al., 2023; Anthropic, 2022). Our analysis reveals that LMs are reluctant to share model uncertainties, de-

spite errors in their generations. LMs can be explicitly prompted to use epistemic markers, but are more likely to generate expressions of certainty than uncertainty, despite an average 47% error rate among high confidence responses, (e.g., *I'm confident the capital of Tanzania is Dar es Salaam.* [Incorrect]).

We then investigate the behavioral responses of users towards model-generated epistemic markers (§4). While linguists and psychologists have long focused on the interpretation of epistemic markers by humans (Budescu et al., 1988; Windschitl and Wells, 1996; Wallsten et al., 1986), the pragmatic implications of the speakers of such markers being AI systems, combined with the known human overreliance on AI (Bussone et al., 2015; Jacobs et al., 2021; Bansal et al., 2021b; Bućinca et al., 2021), could drastically change their interpretation compared to human-spoken ones. Thus, we conduct several user studies to measure how individuals interpret and respond to uncertainties articulated by LMs in calibrated and miscalibrated settings (Figure 1). Our findings surprisingly indicate that users are heavily reliant on LM generated expressions of marked *"I'm sure it's..."* and unmarked certainty (e.g., *"The answer is..."*). Subsequent experiments show that even minor miscalibrations in how a model uses epistemic markers can lead to long-term harms in human performance.

Lastly, given our findings on model overconfidence and human reliability of LM generations, we pinpoint the origins of model overconfidence (§5). We investigate model artifacts such as base models, instruction-tuned models, reward models, and human feedback datasets to isolate the origins of model confidence. Our investigation identifies the process of aligning models with human feedback (i.e., RLHF) as a key contributing factor and we uncover that human annotators are biased **against** expressions of uncertainty.

Together, our findings expose the shortcomings of how LMs currently use epistemic markers, outlines the risks that they pose on downstream users, and put forward mitigating solutions.

2 Epistemic Markers in Language Models

Our work focuses on the alignment between LM accuracy and LM-articulated *epistemic markers* as perceived by users. This is referred to as *linguistic calibration* (Mielke et al., 2020) which builds off of work in both linguistics and machine learning.

Linguists have extensively studied epistemic markers as ubiquitous linguistic features that signal speaker commitment and stance. These markers broadly fall into two categories: **weakeners**—expressions of uncertainty, and **strengtheners**—expressions of certainty (Lakoff, 1975; Hyland, 2005, 2014).

In machine learning, work has focused on improving model calibration (Jiang et al., 2021; Desai and Durrett, 2020; Jagannatha and Yu, 2020; Kamath et al., 2020; Kong et al., 2020) by calibrating the confidence value assigned by a model and model accuracy through a measure called ECE (Naeini et al., 2015). Recent work has focused explicitly on how pretraining (Hendrycks et al., 2019) and scaling (Srivastava et al., 2022; Chen et al., 2023) impacts the calibration of language models. Most relevant to our work is Dhuliawala et al. (2023)'s studies on how humans interpret numerical confidences in calibrated and miscalibrated settings. A key issue remains, as numeric confidence values are known to be challenging for users to interpret (Miller, 2019). Our work, in contrast, aims for a more comprehensive understanding of how humans interpret LM-generated verbal epistemic markers.

We begin by eliciting open-ended generations, a departure from prior methods which prompts models to produce a predefined set of confidence expressions either numerically (Kadavath et al., 2022; Tian et al., 2023; Liu et al., 2023; Xiong et al., 2024; Tanneru et al., 2023) or using an ordinal scale (Lin et al., 2022; Mielke et al., 2020). Instead, our strategy enables a qualitative *bottom-up* (Charmaz, 2006) approach towards understanding how LMs generate epistemic markers, mimicking how real-world users might engage with LMs.

3 How do LMs use Epistemic Markers?

To answer our first motivating question, we investigate how LMs such as GPT, LLaMA-2, and Claude express uncertainties within a broad and challenging, question-answering context. We find that LMs prefer to respond with answers free of epistemic markers and when LMs do use epistemic modifiers, they rely too much on strengtheners, leading to overconfident but incorrect generations.

3.1 Methods

Our objective is to assess the potential harms and safety risks associated with widely used publicly

The following is a multiple choice question about computer security. When replying to the question, incorporate a hedge or a strengthener to signify your certainty level.	SHA-1 has a message digest of (A) 160 bits (B) 512 bits (C) 628 bits (D) 820 bits
	Answer:

Figure 2: Example of the prompt which uses an MMLU question and an instruction which elicits epistemic markers. The **green text** is the category of the question, the **purple text** represents one of the 49 prompts we’ve curated, the **yellow text** is the question and the dark grey text are the multiple choice options.

deployed models like GPT, LLaMA-2, and Claude.¹ We pose a diverse set of questions from the Massive Multitask Language Understanding benchmark (MMLU Hendrycks et al., 2021), a four-way multiple-choice dataset spanning 57 subjects that assess both language model knowledge and problem-solving skills. We design confidence eliciting prompts and measure systematic trends in how LMs use strengtheners and weakeners.

Prompt Design We systematically prompt LMs for epistemic markers by designing three types of open-ended prompt instructions. We modify the *base template* from the original MMLU paper by appending additional instructions crafted to elicit: 1) epistemic markers “Please answer the question and provide your certainty level” (Epi-M), 2) chain-of-thought reasoning (CoT), “Explain your thought process step by step” or 3) a combination of both “Using expressions of uncertainty, explain your thought process step by step” (Epi-M+CoT) (Figure 2). Previous studies have shown that chain-of-thought prompts can enhance model behavior through step-by-step reasoning, and we hypothesized that the process of articulating reasoning might also generate epistemic markers (Wei et al., 2022; Suzgun et al., 2023; Wang et al., 2022).

To ensure the generalizability of our results, we employ snowball sampling to generate a list of diverse prompts, gathering additional paraphrases of prompts from Amazon Mechanical Turk Workers and GPT-3.5 (details in §A).

We prompt nine models (text-davinci-003, GPT-3.5-Turbo, GPT-4, LLaMA-2 7B, LLaMA-2 13B, LLaMA-2 70B, Claude-1, Claude-2, Claude-Instant-1) using 49 prompts on 284 questions, resulting in a total of 125,244 queries.² Using a zero-

¹Models were accessed June - November 2023.

²Duplicated question in the development set of MMLU benchmark, resulting in 284 instead of 285 questions.

	Base	CoT	Epi-M	Epi-M +CoT	Avg*
% in responses	n=1	n=8	n=24	n=16	n=48
all epi. markers	6%	16%	71%	57%	57%
strengtheners	0%	3%	24%	24%	20%
weakeners	2%	3%	15%	20%	14%

Table 1: Models struggle to generate epistemic markers without explicit prompting for strengtheners and weakeners. LMs generate responses with strengtheners despite low accuracies. *Average is across all models and all templates except the base template.

shot prompting approach, we aim to simulate the interactions of end-users. We set the temperature to 0.3 (OpenAI, 2023; Anil et al., 2023), with no stop tokens, and limit the token generation length to a maximum of 400 tokens (details in §C).

Eliciting and Classifying Epistemic Markers

The authors then qualitatively code (Auerbach and Silverstein, 2003; Charmaz, 2006) for epistemic markers in generated responses from each model family via Regex pattern matching and categorizing them as strengtheners and weakeners.³

Specifically, the process involved authors 1) manually looking through the generated responses, 2) identifying codes, 3) designing Regex heuristics to automatically detect markers, and 4) evaluating markers manually. Each evaluation consisted of randomly sampled 100 codes and evaluating them by hand. The entire qualitative coding process was repeated six times until we reached over 90% accuracy. Although this process involved the analysis of over a dozen models, as with most supervised methods, future researchers may need to adapt our schema when using it for out-of-distribution content. Our qualitative coding yielded a total of 76 strengtheners and 105 weakeners. See Tables 7 and 8 for the most commonly generated expressions.

3.2 Findings

Models are reluctant to reveal uncertainties, but can be encouraged.

LMs fail to incorporate epistemic markers in their responses when prompted with the base template. Only 5% of the generated answers include any type of epistemic markers (Table 1), indicating that the majority of responses seen by end-users lack information regarding model uncertainties. We refer to responses that

³A future iteration could include training a classifier to identify codes at the trade-off of lowered transparency and interpretability.

don't contain any epistemic markers as *plain statements* (e.g., “(A)” or “*The answer is (A)*”).

When using chain-of-thought instructions or explicit instructions to elicit epistemic markers, LMs can be encouraged to produce more epistemic markers. Resulting in 16% and 65%⁴ of generations incorporating epistemic markers respectively.

Models are biased towards using strengtheners.

Six out of nine models have a preference to generate significantly more strengtheners than weakeners. Our results indicate that an average of 20% of generations had strengtheners while only 14% had weakeners. This bias is true among prompts eliciting for certainty, with or without CoT (Table 1). We see this trend emerge strongly among GPT and LLaMA-2 chat models, with the Claude-2 being more balanced in its generation of strengtheners and weakeners (Figure 3). Interestingly, the smaller models (LLaMA-2-7B and Claude-Instant-v1⁵) have a higher use of weakeners over strengtheners; model size and generation of epistemic markers could be explored in future work.

Overconfidence results in confident but inaccurate generations. Across all generations, only 53% of generations with expressions of certainty are correct (random accuracy being 25%). Although this accuracy rate is higher than accuracies among weakeners (32%), this is an alarmingly high rate of errors among strengtheners. Furthermore, due to the high rate of generations of strengtheners, 17% of all incorrect answers include strengtheners.

3.3 Discussion

Our findings illustrate that models struggle to appropriately use epistemic markers. First, models are reluctant to produce uncertainties, even when asked with CoT prompts, presenting a veil of tacit certainty. When explicitly prompted to verbalize model confidences, LMs are prone to overuse expressions of certainty even when the output is incorrect, creating potential downstream harms (see §4). The overuse of expressions of certainty is likely to contribute to the existing problem of human overreliance on AI predictions (Jacobs et al., 2021; Bussoni et al., 2015) and explanations (Bansal et al., 2021b; Poursabzi-Sangdeh et al., 2018; Wang and

⁴Weighted average emittance of epistemic markers when prompted with epistemic instructions and epistemic instructions with chain-of-thought

⁵Announced as a “lighter” version of Claude. See: <https://www.anthropic.com/index/introducing-claude>




Figure 3: Use of strengtheners and weakeners in generations across GPT, LLaMA-2, and Claude Models. Confidence intervals calculated with bootstrap resampling.

Yin, 2021; Ehsan et al., 2021).

The linguistic miscalibration is emerging as a new safety risk as LMs play a bigger role in human-LM collaborations. In the space of information seeking, the need for uncertainties and knowledge limitations will likely continue to persist even as model performance improves. Moving forward, we must examine how to mitigate the harms of model overconfidence and how to best build cognitive forcing designs, such as verbalized uncertainties, to discourage human overreliance of AI systems (Buçinca et al., 2021).

4 Human Interpretations of Uncertainty

With more robust understanding of how LMs use epistemic markers, we shift our focus to the second inquiry: how do humans interpret LM-generated epistemic markers? Using a subset of expressions generated by LMs, we set up a task to evaluate the effect of these markers on user reliance on AI. We find that users by default are highly reliant on LM-generated responses and that even minor miscalibrations in systems can have long-term consequences in human-LM collaborations.

4.1 Methods

Creating a Self-Incentivized Task We create a self-incentivized task where users must accrue points by correctly answering challenging trivia questions, deciding whether to rely on an AI agent’s response for help. We situate users in an imagined game scenario where they are asked to interact with AI agent named Marvin, a set up adopted from an user-AI interaction study from




Figure 4: Example of Setting 1 human experiments task.

Bansal et al. (2019a).⁶ In the game, the user shown a question (e.g., "What is the capital of Palau?") and a response generated by Marvin that includes epistemic markers (e.g., "I'm certain its Ngerulmud"). The user must decide whether to rely on Marvin's answer or whether to indicate that they'll look it up themselves later.

Trivia Question Selection We control for uniformity in scenario content by limiting questions to country capital trivia. Our task should ensure that the users' assessment of AI reliability stems from the model's use of epistemic cues rather than users' own prior knowledge. Hence, we select the most challenging trivia questions as ranked from Sporcle, an online trivia platform.⁷ By selecting countries where participants are unlikely to know the answer, we encourage users to primarily use LM-generated epistemic markers rather than their own knowledge to make decisions.

Recruitment Process We launch the task using Prolific and Qualtrics and inform the participants of the nature and the risks of the task through a consent form. The task is compliant with internal review board (IRB) protocols.

Template Selection We select the most frequently occurring expressions of certainty and uncertainty from §3 and transform them into prefixes in the context of question answering (e.g., "I think it's", "Perhaps it's"). We filter for naturalistic expressions, avoiding any template duplication.

For details on recruitment process, IRB protocol, and template selection see Appendix B.

4.2 Experimental Settings

Setting 1: Control Setting We first use our task to evaluate how participants rely on LM generated epistemic markers (from §3). Participants are

shown a set of trivia questions and *the beginning of a response* (e.g., "I think it's..." Figure 4). Users are asked whether or not they'd like to rely on Marvin's answer. Since the users do not see any answers, they are simply expressing their reliance of epistemic markers as generated by LMs. Participants are presented with strengtheners, weakeners, and plain statements (expressions free of epistemic markers like, "The answer is (A)"). We recruited 25 participants and each were shown 106 questions.

Setting 2: Interactive Settings The next three settings are interactive settings. Participants engage in 50 rounds of question-answering where in each round they are 1) shown a question, 2) shown Marvin's predicted response with epistemic markers, 3) asked to make a decision, and 4) given feedback on their decision. Providing users with feedback gives users the opportunity to build a mental model of how Marvin performs (Bansal et al., 2019a) and allows us to measure the harms that may arise from long-term interaction (Lee et al., 2022). We recruited 25 new participants for each setting.

In these experimental settings, we introduce a scoring system, also modified from Bansal et al. 2019a. The scoring set-up is designed such that the only way to have a positive score is to rely on Marvin correctly based on relying the epistemic markers (see Table 2).⁸

User Action	Marvin Correct	Marvin Incorrect
Rely on Marvin	1	-1
Look up answer	0	0

Table 2: Scoring chart for human experiments. Relying on Marvin's generation when Marvin is right will yield 1 point, -1 otherwise. Looking up the answer yields 0 points. Half the answers are wrong, so choosing to always to rely on Marvin or to look up the answer will yield a total score of 0.

Setting 2A: Calibrated Setting In the calibrated setting, Marvin's responses are calibrated with the expected human interpretations of epistemic markers from the control setting (i.e., strengtheners appear with correct answers and weakeners appear with incorrect answers).

Setting 2B: Overconfident Setting In the overconfident setting, Marvin will use strengtheners

⁶The original task involves users classifying shapes. We adopt the decision making setup from this work.

⁷<https://www.sporcle.com/games/g/worldcapitals/results>

⁸To avoid spammers, we filtered the bottom 20% participants based on their performance on this task.




Figure 5: Participant results in the calibrated, overconfident, and underconfident settings. We see lower scores across the miscalibrated rounds. In the overconfident setting, lowered scores persist in the later calibrated rounds as well.

when generating the incorrect answer (e.g., “I’m sure the capital of Vanuatu is Luganville” [Incorrect]). This setting has five overconfident responses which will appear in the first 20 questions; the remaining 30 questions are calibrated.⁹

Setting 2C: Underconfident Setting In the underconfident setting, Marvin will use weakeners when generating the correct answer (e.g., “I’m not certain, but maybe the capital of Vanuatu is Port Vila” [Correct]). Again, five underconfident responses will appear in the first twenty questions.

4.3 Findings

Users rely on strengtheners but also on plain statements In the control setting, when presented with expressions of strengtheners, nearly 90% of users will rely on Marvin’s response. When presented with weakeners, approximately 90% of users will choose to look up the answer (see Table 6 for details). Surprisingly, plain statements like “The answer is (A)” or “(A)” are also relied on by users nearly 90% of the time. In other words, without communicating any epistemic markers, humans interpret this as a sign of model certainty.

Users Effectively Leverage Calibrated LM-Generated Epistemic Markers In the interactive calibrated setting, our results illustrate that users are able to learn mental models of epistemic markers after approximately 20 rounds. In the early rounds, users rely on strengtheners 94% of the time and on weakeners 7% of the time. After 20 rounds, most participants are able to nearly perfectly leverage Marvin’s epistemic markers with users learning to rely on strengtheners 99% of the time and weakeners nearly 1% of the time, averaging an accuracy

of 97% across all rounds. The high performance in this calibrated setting also validates that participants are primarily relying on epistemic markers, to answer these questions. If they had been only relying on their own knowledge, it would be unlikely to see participants perform nearly perfectly on such challenging trivia questions.

Users are Overreliant on Overconfident Responses In the first 20 rounds, users relied on 81% of strengtheners, when only 66% of the generations with strengtheners were correct, accruing on greater penalties than necessary. Because participants are unlikely to know the answer to the question, we see that participants mistakenly rely on incorrect, confident generations an average 73% of the time. The consequences of system miscalibration continued to negatively impact human performance, even after several rounds of calibrated model answers. Participants averaged 76% in the miscalibrated rounds and 86% on the calibrated rounds (Figure 5).

Users in the Overconfident Setting Also Incorrectly Relied on Weakeners An unexpected effect of the overconfident responses was that participants started to interpret weakeners differently. Even though weakeners were used correctly in this setting (i.e., none of the answers with weakeners were correct), the participants were observed to rely on answers with weakeners at a higher rate than in the control setting (9% vs. 3%).

Users are Underreliant on Weakeners in Underconfident LMs Participants in the underconfident setting averaged 66% in the miscalibrated rounds but 98% on the later calibrated rounds, matching the performance in the calibrated setting. This is in contrast to the overconfident setting where users’ mental models were never fully cor-

⁹To create consistency in how the answers are incorrect, the largest non-capital city is used instead of the capital city.

rected, even after the same number of calibrated rounds.

4.4 Discussion

Plain Statements are *Confident* Statements

The troubling concern that current models struggle to use epistemic markers in calibrated ways (§3) is underscored by our finding that lack of epistemic markers is perceived as confident language (§4). That is, the absence of calibrated markers presents significant harms for human-AI collaboration. For example, LM hallucinations are not only factually incorrect (Ji et al., 2022; Maynez et al., 2020; Zhang et al., 2023), but as our study suggests, they may also be interpreted as high certainty due to the lack of epistemic markers. One potential design recommendation is to generate weakeners without explicit elicitation and use plain statements only when the model is confident.

Miscalibrations in Strengtheners Impact Interpretation of Weakeners

Miscalibration in the use of strengtheners resulted in users interpreting weakeners incorrectly as well. Our findings signal that miscalibration in one dimensional (e.g., the incorrect use of strengtheners leads users to distrust other uses of epistemic markers). As a recommendation, researchers measuring the harms of miscalibration must also consider how miscalibration impacts the users' mental models of the whole system, rather than the perception of just an individual component. This work ties into existing literature on how mental models of AI systems are affected by numerous complexities such as accuracy (Bansal et al., 2021a), warmth (McKee et al., 2024), and system updates (Bansal et al., 2019b).

Long-Term Effects of Overconfidence Our findings signal that mental models of language models are developed early in LM-interactions, potentially resulting in long-term harms, even after models become calibrated later on. This corroborates work from Dhuliawala et al. (2023) who conducted human interpretations of numerical uncertainties from LMs. Similarly, we find that users can correct a mental model developed from an underconfident model but struggle to do so with overconfident models. Unfortunately, public models are overconfident, creating not only a reliability harm now, but also potentially creating long-term algorithmic aversion (Dietvorst et al., 2015) to future models.

5 Origin of Model Overconfidence

Our work shows that LMs are overconfident and that humans are highly reliant on plain statements and statements with strengtheners; all of which is exacerbated by models often being incorrect when expressing certainty. Here, we turn to our last motivating question: What is the origin of model confidence and what are potential mitigations? In this section, we pinpoint LM overconfidence to an artifact of the post-training alignment process, specifically a bias from human annotators against uncertainty.

5.1 Where Does LM Overconfidence Begin?

Current state-of-the-art models are trained using a number of techniques to support human-AI collaboration through natural language. Starting with a pretrained *base* model, one of the most popular techniques is to use *supervised fine-tune* (SFT) using human instructions. The model is then trained with *reinforcement learning with human feedback* (RLHF), where a reward model is learned through human preferences of pairwise text comparisons.

5.2 Methods

Model Stages We perform analysis on the models from the GPT and LLaMA-2 family to identify the origin of model overconfidence. Using the same prompting strategy as §3, we measure how base models and supervised fine-tuned models compare to their RLHF counterparts when it comes to generating expressions of certainty.

We compare three models from the GPT3 family, *davinci*, *text-davinci-002*, and *text-davinci-003* which are base, supervised fine-tuned, and RLHF models, respectively.¹⁰ We then compare LLaMA-2 models base models with their SFT+RLHF counterparts.

Reward Modeling We directly probe OpenAssistant's open-sourced reward model trained on human feedback datasets and assess their scoring.¹¹ We prompt the model with a question-response pair where the question is "What is the capital of X?" and the response is an epistemic marker like "*I think it's*". We test the reward model on 183 question answer pairs across a subset of 30 commonly occurring templates generated by LMs in §3 and

¹⁰Models were accessed June - November 2023.

¹¹<https://huggingface.co/OpenAssistant/reward-model-deberta-v3-large-v2>




Figure 6: Base models vs. RLHF models in their generation of strengtheners and weakeners. In base models, we see a preference for weakeners but the trend reverses among RLHF models.

compare the model scores with human judgements from §4.

Human Annotated Datasets Lastly, we examine the datasets that were used to train the open-sourced reward model. Specifically, we examine the datasets: OpenAI’s “WebGPT comparison” and “Summarize with Feedback”, Dahoa’s “Synthetic Instruct GPT Pairwise” dataset, and Anthropic’s “Helpful and Harmless” dataset.¹¹ We then measure how often strengtheners and weakeners are preferred by human annotators in these datasets.

5.3 Findings

Overconfidence in RLHF Models We quantitatively observe that RLHF-ed models emit more strengtheners than weakeners, which contrasts to the base and instruction-tuned variants where the pattern is the opposite (Figure 6). This suggests that this preference for strengtheners is introduced during the RLHF process.

Reward Modeling Is Biased Towards Certainty

Reward modeling prefers plain statements with an average score of 4.03, followed by strengtheners with a score of 0.82. However, there is a strong penalty applied to weakeners, with the average rewards score of -1.86. Lower scores in reward modeling would in turn influence how LMs generate natural language, leading to a bias in avoiding the language of uncertainty at generation time. See Table 3 for a comparison between reward model and human scoring.

Human Raters are Biased Against Uncertainty

We annotate texts in each datasets as containing strengtheners and weakeners and measure how these rates vary across the chosen and rejected texts as annotated by human judges. Given the prevalence of model overconfidence, we hypothesized

that there may be a preference for strengtheners in chosen texts. We find that this is not the case. In fact, there is a slight, but significant preference for strengtheners in rejected texts as compared to chosen text (2.95% vs. 2.72%).¹² In a pairwise comparison, we see that plain text is actually slightly preferred over strengthened texts (chosen 9% more often). This shows that in these datasets, there is not a human bias for strengtheners. Weakeners also appears significantly more often among rejected texts (5.02%) compared to chosen texts (4.47%). In a pairwise comparison, weakeners are preferred 8% less often than strengtheners and 9% less often than plain texts. This highlights that annotators don’t have a bias for certainty, rather there is a bias *against* weakeners, matching the results seen from the reward model experiments.

5.4 Discussion

Uncovering Unknowns in Human Preferences

Feedback Alignment Our investigation into RLHF datasets reveals that humans have implicit biases towards other dimensions of language which may not be known in the annotation phase. Although our study focuses RLHF, the process of aligning language models to biased human feedback will remain as issue, regardless of the exact training method (e.g., DPO or SLiC) (Zhao et al., 2023; Rafailov et al., 2024). The artifact of humans having a bias against uncertainty adds to the list of implicit biases in annotations (Gururangan et al., 2018) (e.g., text length (Saito et al., 2023), toxicity detection (Sap et al., 2022b), political leanings (Santurkar et al., 2023)). However, the human bias against uncertain language is particularly harmful as it causes aligned models to be reluctant in their generation uncertainty, negatively impacting human over-reliance on LMs. Interventions such as swapping labels on annotated datasets could potentially mitigate some overconfidence harms but could also risk introducing new, yet to be known biases. Investigations are needed in human feedback datasets and annotation processes to fully understand the potential implicit biases that are introduced through human preference annotations.

Beyond Mimicking Human Language

As LMs evolve towards generating more nuanced language, such as uncertainties, a shift in design is prudent. The work of Hollan and Stornetta (1992) discusses the need to design technology that goes beyond

¹²95% CI calculated using bootstrap sampling.

simply mimicking that of the real world. We could apply this design thinking towards designing natural language as an interface. Instead of building language models simply based on mimicking human preferences, we could instead design LMs to verbalize uncertainty in ways that would increase cognitive engagement and lower human overreliance (Bućinca et al., 2021).

6 Desired Criteria and Mitigating Solutions Moving Forward

Our research highlights the overconfidence of language models, the reliance of humans on generations, and the long-term consequences of miscalibration in human-AI interactions. Here, we propose three criteria and potential implementations to mitigate overconfidence in LMs.

Criteria: Unsolicited Epistemic Markers Language models should autonomously emit expressions of uncertainty without prompting, akin to human behavior, in order to appropriately convey levels of confidence. This is critical as the lack of epistemic marker emittance is perceived by humans as tacit certainty.

Concretely, this could be addressed in various stages of the RLHF pipeline, including **data augmentation, annotator training, or dataset evaluation**. RLHF datasets could be augmented to include additional and more representative epistemic markers. Similar studies have been successful in improving safety (Ji et al., 2023), modeling different perspectives (Dong et al., 2024), and performing moral self-correction (Ganguli et al., 2023). Annotators could also receive training (Clark et al., 2021) and priming (Sap et al., 2019) to be cognizant of potential biases (e.g., biases against uncertainty). Lastly, automatic sensibility checks and data filtering using our heuristics could be used to evaluate the effectiveness of the proceeding interventions, building transparency towards the levels of overconfidence in RLHF datasets.

Criteria #2: Comprehensive Coverage Language models ought to have the capacity to generate and interpret the full range of epistemic markers. However, current pre-training data is often limited to written internet data, leading to numerous limitations and biases (Gordon and Van Durme, 2013; Lucy and Gauthier, 2017; Sap et al., 2022a).

Concretely, incorporating **more diverse data sources** would allow models to learn a larger range

of epistemic markers. For example, hedges have long been studied in speech in formal and informal settings (Aijmer, 1986, 2013) such as presidential debates and speeches (Al-Rashady, 2012; Mansour and Alghazo, 2021), TED talks (Nuraniwati and Permatasari, 2021), and student presentations (Muziatun et al., 2021). Leveraging alternative training sources such as podcasts, news transcripts, and peer conversations could broaden the use of expressions of uncertainty by LMs.

Criteria #3: Context-Dependent Calibration

In addition to being internally calibrated (i.e., calibrated between internal probabilities and verbalized certainties), LMs should also be context-dependent in their calibration. If a model were deployed for entertainment purposes, a lower threshold for expressions of certainty could be tolerated; but if the same language model were to be deployed to a mission-critical task, new thresholds based on *the context* must be determined.

Concretely, this could be accomplished via **user calibration or system-level prompts**. For example, one could implement an initialization procedure, similar to the procedure shown in §4.2, that allows practitioners to measure how humans would rely on LM-generated expressions in a specific context. The results of this procedure would then inform practitioners and users of how to best to fit the model to the specific needs using few-shot prompting, fine-tuning, or system-level prompts. Future work could include investigate online algorithms that adjust the certainty thresholds based on the human-LM initialize round.

7 Conclusion

Our work set out to explore how users interpret epistemic markers as generated by LMs in an effort to better understand the shortcomings of human-LM communications. We find that LMs are overconfident in their generations and that users are highly reliant on LM responses whether there is implicit or explicit confidence. We trace the origin of model overconfidence to the RLHF process and find that annotators have a bias against uncertainty in text.

8 Limitations

Cultural Interpretations of Epistemic Markers

Our study focused exclusively on how language models generate epistemic markers in the English language (Bender, 2019). Humans greatly differ

in their use of hedges and strengtheners across languages, contexts, and cultures (Itani, 1995; Lauwereyns, 2002; Yagız and Demir, 2014; Nguyen Thi Thuy, 2018; Mur-Dueñas, 2021), future studies could consider how non-English language models differ in their use of strengtheners and weakeners.

Our human studies recruited U.S. based participants exclusively and their willingness to rely on epistemic markers is shaped by their experiences and cultural context. Our results thus illustrate a narrow and U.S.-centric view of how humans might interpret epistemic markers (Henrich et al., 2010; Atari et al., 2023). Participants from other cultural backgrounds might reveal different findings on how humans rely on LM-generated epistemic markers.

The Ambiguity of Weakened Strengtheners

LMs articulated epistemic markers that fell in between strengtheners and weakeners, which we labeled as labeled as *weakened-strengtheners*. This schema closely follows Sanders and Spooren (1996)’s schema of *certainty*, *uncertainty*, and *semi-certainty* markers. In our human experiments, we found that participants displayed great variance in their reliance of weakened strengtheners as some humans appear to rely on weakened strengtheners meanwhile others do not (see Table 6). This ambiguous interpretation of weakened strengtheners highlights a potential risk for miscommunication; the prolific use of them could lead to more confusion and misinterpretation than clarity. Further work on the more nuanced features of epistemic markers is needed.

Gap Between Human Experiments and Real Self-Incentivized Users

A gap still exists between self-incentivized users and the participants who we recruited for our experiments. Despite our best efforts to situate users in a real-life scenario, the harms we uncover here will likely differ from that of users in a real deployed setting. The contexts in which users might engage with these chat models would like also influence their interpretability of epistemic markers (Grice, 1975; Goodman and Frank, 2016). Changes such as the metaphors associated with the agent itself could have significant impacts on user reliance behaviors (Khadpe et al., 2020). Further investigations and in-depth user studies and interviews may be needed to comprehensively study the harms of LM overconfidence.

9 Ethics Statement

Our paper is primarily considered with the potential harms and ethical implications that may arise when humans interact with LLMs. Our findings illustrate that LLMs systematically fail to represent model uncertainties, creating a threat that results humans overrelying on incorrect generations. In regards to our human experiments, we followed standard practices such as providing participants with informed consent, paying participants, on average, over \$15 USD/per hour and debriefing participants on the correct answers if they made errors when interacting with our system.

Acknowledgements

Thank you so much to Dan Jurafsky, Yejin Choi, Myra Cheng, Kristina Gligorić, Abhilasha Ravichander, Tal August, Luca Soldaini and the AI2 Mosaic team for their helpful feedback and advice! Thank you so much to the numerous participants who helped us with our pilot and formal human studies. Kaitlyn Zhou is supported by the Stanford Graduate Fellowship. Xiang Ren’s research is supported in part by the Office of the Director of National Intelligence (ODNI), Intelligence Advanced Research Projects Activity (IARPA), via the HIATUS Program contract #2022-22072200006, the Defense Advanced Research Projects Agency with award HR00112220046, and NSF IIS 2048211.

References

- Karin Aijmer. 1986. Discourse variation and hedging. In *Corpus linguistics II*, pages 1–18. Brill.
- Karin Aijmer. 2013. *Understanding pragmatic markers: A variational pragmatic approach*. Edinburgh University Press.
- Fahad Al-Rashady. 2012. Determining the role of hedging devices in the political discourse of two american presidentiables in 2008. *TESOL Journal*, 7(1):30–42.
- Rohan Anil, Andrew M Dai, Orhan Firat, Melvin Johnson, Dmitry Lepikhin, Alexandre Passos, Siamak Shakeri, Emanuel Taropa, Paige Bailey, Zhifeng Chen, et al. 2023. *PaLM 2 technical report*. *arXiv preprint arXiv:2305.10403*.
- Anthropic. 2022. *Introducing claude*.
- Mohammad Atari, Mona J. Xue, Peter S. Park, Damián E. Blasi, and Joseph Henrich. 2023. *Which humans?*

- Carl F. Auerbach and Louise Bordeaux Silverstein. 2003. [Qualitative data: An introduction to coding and analysis](#).
- Austin S Babrow, Chris R Kasch, and Leigh A Ford. 1998. [The many meanings of uncertainty in illness: Toward a systematic accounting](#). *Health communication*, 10(1):1–23.
- Gagan Bansal, Besmira Nushi, Ece Kamar, Eric Horvitz, and Daniel S Weld. 2021a. [Is the Most Accurate AI the Best Teammate? Optimizing AI for Teamwork](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 11405–11414.
- Gagan Bansal, Besmira Nushi, Ece Kamar, Walter S Lasecki, Daniel S Weld, and Eric Horvitz. 2019a. [Beyond Accuracy: The Role of Mental Models in Human-AI Team Performance](#). In *Proceedings of the AAAI conference on human computation and crowdsourcing*, volume 7, pages 2–11.
- Gagan Bansal, Besmira Nushi, Ece Kamar, Daniel S Weld, Walter S Lasecki, and Eric Horvitz. 2019b. [Updates in Human-AI Teams: Understanding and Addressing the Performance/Compatibility Tradeoff](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 2429–2437.
- Gagan Bansal, Tongshuang Wu, Joyce Zhou, Raymond Fok, Besmira Nushi, Ece Kamar, Marco Tulio Ribeiro, and Daniel Weld. 2021b. [Does the whole exceed its parts? the effect of ai explanations on complementary team performance](#). In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, CHI '21, New York, NY, USA. Association for Computing Machinery.
- Emily Bender. 2019. [The#benderrule: On naming the languages we study and why it matters](#). *The Gradient*, 14.
- Rishi Bommasani, Drew A Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, et al. 2021. [On the opportunities and risks of foundation models](#). *arXiv preprint arXiv:2108.07258*.
- Dale E Brashers, Judith L Neidig, Stephen M Haas, Linda K Dobbs, Linda W Cardillo, and Jane A Russell. 2000. [Communication in the management of uncertainty: The case of persons living with hiv or aids](#). *Communications Monographs*, 67(1):63–84.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. [Language Models are Few-Shot Learners](#). *Advances in neural information processing systems*, 33:1877–1901.
- Zana Bućinca, Maja Barbara Malaya, and Krzysztof Z. Gajos. 2021. [To trust or to think: Cognitive forcing functions can reduce overreliance on ai in ai-assisted decision-making](#). *Proc. ACM Hum.-Comput. Interact.*, 5(CSCW1).
- David V Budescu, Shalva Weinberg, and Thomas S Wallsten. 1988. [Decisions based on numerically and verbally expressed uncertainties](#). *Journal of Experimental Psychology: Human Perception and Performance*, 14(2):281.
- Adrian Bussone, Simone Stumpf, and Dymna O’Sullivan. 2015. [The role of explanations on trust and reliance in clinical decision support systems](#). In *2015 International Conference on Healthcare Informatics*, pages 160–169.
- Carrie J. Cai, Emily Reif, Narayan Hegde, Jason D. Hipp, Been Kim, Daniel Smilkov, Martin Wattenberg, Fernanda B. Viégas, Gregory S. Corrado, Martin C. Stumpe, and Michael Terry. 2019. [Human-Centered Tools for Coping with Imperfect Algorithms During Medical Decision-Making](#). *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*.
- Kathy Charmaz. 2006. *Constructing grounded theory: A practical guide through qualitative analysis*. SAGE.
- Yangyi Chen, Lifan Yuan, Ganqu Cui, Zhiyuan Liu, and Heng Ji. 2023. [A close look into the calibration of pre-trained language models](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1343–1367, Toronto, Canada. Association for Computational Linguistics.
- Elizabeth Clark, Tal August, Sofia Serrano, Nikita Haduong, Suchin Gururangan, and Noah A. Smith. 2021. [All that’s ‘human’ is not gold: Evaluating human evaluation of generated text](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7282–7296, Online. Association for Computational Linguistics.
- Maria De-Arteaga, Riccardo Fogliato, and Alexandra Chouldechova. 2020. [A case for humans-in-the-loop: Decisions in the presence of erroneous algorithmic scores](#). In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*, pages 1–12.
- Shrey Desai and Greg Durrett. 2020. [Calibration of pre-trained transformers](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 295–302, Online. Association for Computational Linguistics.
- Shehzaad Dhuliawala, Vilém Zouhar, Mennatallah El-Assady, and Mrinmaya Sachan. 2023. [A diachronic perspective on user trust in AI under uncertainty](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5567–5580, Singapore. Association for Computational Linguistics.

- Berkeley J Dietvorst, Joseph P Simmons, and Cade Massey. 2015. [Algorithm Aversion: People Erroneously Avoid Algorithms after Seeing Them Err](#). *Journal of Experimental Psychology: General*, 144(1):114.
- Zibin Dong, Yifu Yuan, Jianye HAO, Fei Ni, Yao Mu, YAN ZHENG, Yujing Hu, Tangjie Lv, Changjie Fan, and Zhipeng Hu. 2024. [Aligndiff: Aligning diverse human preferences via behavior-customisable diffusion model](#). In *The Twelfth International Conference on Learning Representations*.
- Marek J Druzdzel. 1989. [Verbal uncertainty expressions: Literature review](#). Pittsburgh, PA: Carnegie Mellon University, Department of Engineering and Public Policy, pages 1–13.
- Upol Ehsan, Q. Vera Liao, Michael Muller, Mark O. Riedl, and Justin D. Weisz. 2021. [Expanding explainability: Towards social transparency in ai systems](#). CHI '21, New York, NY, USA. Association for Computing Machinery.
- Deep Ganguli, Amanda Askell, Nicholas Schiefer, Thomas I Liao, Kamilė Lukošiuūtė, Anna Chen, Anna Goldie, Azalia Mirhoseini, Catherine Olsson, Danny Hernandez, et al. 2023. [The capacity for moral self-correction in large language models](#). *arXiv preprint arXiv:2302.07459*.
- Noah D. Goodman and Michael C. Frank. 2016. [Pragmatic Language Interpretation as Probabilistic Inference](#). *Trends in Cognitive Sciences*, 20(11):818–829.
- Jonathan Gordon and Benjamin Van Durme. 2013. [Reporting bias and knowledge acquisition](#). In *Proceedings of the 2013 workshop on Automated knowledge base construction*, pages 25–30.
- Herbert P Grice. 1975. Logic and conversation. In *Speech acts*, pages 41–58. Brill.
- Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel Bowman, and Noah A. Smith. 2018. [Annotation artifacts in natural language inference data](#). In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pages 107–112, New Orleans, Louisiana. Association for Computational Linguistics.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021. [Measuring massive multitask language understanding](#). *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Dan Hendrycks, Kimin Lee, and Mantas Mazeika. 2019. [Using pre-training can improve model robustness and uncertainty](#). In *International conference on machine learning*, pages 2712–2721. PMLR.
- Joseph Henrich, Steven J. Heine, and Ara Norenzayan. 2010. [Most people are not WEIRD](#). *Nature*, 466:29–29.
- Jim Hollan and Scott Stornetta. 1992. [Beyond being there](#). In *Proceedings of the SIGCHI conference on Human factors in computing systems*, pages 119–125.
- Ken Hyland. 2005. Stance and engagement: A model of interaction in academic discourse. *Discourse studies*, 7(2):173–192.
- Ken Hyland. 2014. Disciplinary discourses: Writer stance in research articles. In *Writing: Texts, processes and practices*, pages 99–121. Routledge.
- Reiko Itani. 1995. *Semantics and pragmatics of hedges in English and Japanese*. University of London, University College London (United Kingdom).
- Maia Jacobs, Melanie F Pradier, Thomas H McCoy Jr, Roy H Perlis, Finale Doshi-Velez, and Krzysztof Z Gajos. 2021. [How machine-learning recommendations influence clinician treatment selections: the example of antidepressant selection](#). *Translational psychiatry*, 11(1):108.
- Abhyuday Jagannatha and Hong Yu. 2020. [Calibrating structured output predictors for natural language processing](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 2078–2092, Online. Association for Computational Linguistics.
- Jiaming Ji, Mickel Liu, Juntao Dai, Xuehai Pan, Chi Zhang, Ce Bian, Ruiyang Sun, Yizhou Wang, and Yaodong Yang. 2023. [Beavertails: Towards improved safety alignment of llm via a human-preference dataset](#). *ArXiv*, abs/2307.04657.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Delong Chen, Wenliang Dai, Andrea Madotto, and Pascale Fung. 2022. [Survey of Hallucination in Natural Language Generation](#). *ACM Computing Surveys*, 55:1 – 38.
- Zhengbao Jiang, Jun Araki, Haibo Ding, and Graham Neubig. 2021. [How can we know when language models know? on the calibration of language models for question answering](#). *Transactions of the Association for Computational Linguistics*, 9:962–977.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield Dodds, Nova DasSarma, Eli Tran-Johnson, et al. 2022. [Language models \(mostly\) know what they know](#). *arXiv preprint arXiv:2207.05221*.
- Amita Kamath, Robin Jia, and Percy Liang. 2020. [Selective question answering under domain shift](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5684–5696, Online. Association for Computational Linguistics.
- Pranav Khadpe, Ranjay Krishna, Li Fei-Fei, Jeffrey T. Hancock, and Michael S. Bernstein. 2020. [Conceptual metaphors impact perceptions of human-ai collaboration](#). *Proc. ACM Hum.-Comput. Interact.*, 4(CSCW2).

- René F. Kizilcec. 2016. [How much information?: Effects of transparency on trust in an algorithmic interface](#). *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*.
- Lingkai Kong, Haoming Jiang, Yuchen Zhuang, Jie Lyu, Tuo Zhao, and Chao Zhang. 2020. [Calibrated language model fine-tuning for in- and out-of-distribution data](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1326–1340, Online. Association for Computational Linguistics.
- George Lakoff. 1975. Hedges: A study in meaning criteria and the logic of fuzzy concepts. In *Contemporary Research in Philosophical Logic and Linguistic Semantics: Proceedings of a Conference Held at the University of Western Ontario, London, Canada*, pages 221–271. Springer.
- Shizuka Lauwereyns. 2002. [Hedges in Japanese conversation: The influence of age, sex, and formality](#). *Language Variation and Change*, 14(2):239–259.
- Mina Lee, Megha Srivastava, Amelia Hardy, John Thickstun, Esin Durmus, Ashwin Paranjape, Ines Gerard-Ursin, Xiang Lisa Li, Faisal Ladhak, Frieda Rong, et al. 2022. [Evaluating human-language model interaction](#). *arXiv preprint arXiv:2212.09746*.
- Stephanie C. Lin, Jacob Hilton, and Owain Evans. 2022. [Teaching models to express their uncertainty in words](#). *Trans. Mach. Learn. Res.*, 2022.
- Kevin Liu, Stephen Casper, Dylan Hadfield-Menell, and Jacob Andreas. 2023. [Cognitive dissonance: Why do language model outputs disagree with internal representations of truthfulness?](#) In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 4791–4797, Singapore. Association for Computational Linguistics.
- Li Lucy and Jon Gauthier. 2017. [Are distributional representations ready for the real world? evaluating word vectors for grounded perceptual meaning](#). In *Proceedings of the First Workshop on Language Grounding for Robotics*, pages 76–85, Vancouver, Canada. Association for Computational Linguistics.
- Enas Mansour and Sharif Alghazo. 2021. [Hedging in political discourse: The case of trump’s speeches](#). *Jordan Journal of Modern Languages and Literatures*.
- Joshua Maynez, Shashi Narayan, Bernd Bohnet, and Ryan McDonald. 2020. [On faithfulness and factuality in abstractive summarization](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1906–1919, Online. Association for Computational Linguistics.
- Kevin R McKee, Xuechunzi Bai, and Susan T Fiske. 2024. [Warmth and competence in human-agent co-operation](#). *Autonomous Agents and Multi-Agent Systems*, 38(1):23.
- Sabrina J. Mielke, Arthur Szlam, Emily Dinan, and Y-Lan Boureau. 2020. [Reducing conversational agents’ overconfidence through linguistic calibration](#). *Transactions of the Association for Computational Linguistics*, 10:857–872.
- Tim Miller. 2019. [Explanation in artificial intelligence: Insights from the social sciences](#). *Artificial intelligence*, 267:1–38.
- Pilar Mur-Dueñas. 2021. [There may be differences: Analysing the use of hedges in english and spanish research articles](#). *Lingua*, 260:103131.
- Muziatun Muziatun, Fahria Malabar, and Nur Sangketa. 2021. [An analysis of hedging devices on students’ presentation of seminar on language based on the gender](#). *Ideas: Jurnal Pendidikan, Sosial, dan Budaya*, 7:97.
- Mahdi Pakdaman Naeini, Gregory F. Cooper, and Milos Hauskrecht. 2015. [Obtaining well calibrated probabilities using bayesian binning](#). *Proceedings of the ... AAAI Conference on Artificial Intelligence. AAAI Conference on Artificial Intelligence*, 2015:2901–2907.
- Thu Nguyen Thi Thuy. 2018. [A corpus-based study on cross-cultural divergence in the use of hedges in academic research articles written by vietnamese and native english-speaking authors](#). *Social Sciences*, 7(4).
- Tri Nuraniwati and Alfelis Nugky Permatasari. 2021. [Hedging in ted talks: A corpus-based pragmatic study](#). *JEELS (Journal of English Education and Linguistics Studies)*, 8(2):203–226.
- OpenAI. 2023. [GPT-4 technical report](#).
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke E. Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Francis Christiano, Jan Leike, and Ryan J. Lowe. 2022. [Training language models to follow instructions with human feedback](#). *ArXiv*, abs/2203.02155.
- Forough Poursabzi-Sangdeh, Daniel G. Goldstein, Jake M. Hofman, Jennifer Wortman Vaughan, and Hanna M. Wallach. 2018. [Manipulating and measuring model interpretability](#). *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. 2024. [Direct preference optimization: Your language model is secretly a reward model](#). *Advances in Neural Information Processing Systems*, 36.
- Keita Saito, Akifumi Wachi, Koki Wataoka, and Youhei Akimoto. 2023. [Verbosity bias in preference labeling by large language models](#). *arXiv preprint arXiv:2310.10076*.

- José Sanders and Wilbert Spooren. 1996. Subjectivity and certainty in epistemic modality: A study of dutch epistemic modifiers.
- Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cinoo Lee, Percy Liang, and Tatsunori Hashimoto. 2023. [Whose opinions do language models reflect?](#) In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 29971–30004. PMLR.
- Maarten Sap, Dallas Card, Saadia Gabriel, Yejin Choi, and Noah A. Smith. 2019. [The risk of racial bias in hate speech detection](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 1668–1678, Florence, Italy. Association for Computational Linguistics.
- Maarten Sap, Ronan Le Bras, Daniel Fried, and Yejin Choi. 2022a. [Neural theory-of-mind? on the limits of social intelligence in large LMs](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3762–3780, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Maarten Sap, Swabha Swayamdipta, Laura Vianna, Xuhui Zhou, Yejin Choi, and Noah A. Smith. 2022b. [Annotators with attitudes: How annotator beliefs and identities bias toxic language detection](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 5884–5906, Seattle, United States. Association for Computational Linguistics.
- Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R. Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, Agnieszka Kluska, Aitor Lewkowycz, Akshat Agarwal, Alethea Power, Alex Ray, Alex Warstadt, Alexander W. Kocurek, Ali Safaya, Ali Tazarv, Alice Xiang, Alicia Parrish, Allen Nie, Aman Hussain, Amanda Askell, Amanda Dsouza, Ambrose Slone, Ameet Annasaheb Rahane, Anantharaman S. Iyer, Anders Andreassen, Andrea Madotto, Andrea Santilli, Andreas Stuhlmüller, Andrew M. Dai, Andrew La, Andrew Kyle Lampinen, Andy Zou, Angela Jiang, Angelica Chen, Anh Vuong, Animesh Gupta, Anna Gottardi, Antonio Norelli, Anu Venkatesh, Arash Gholamidavoodi, Arfa Tabasum, Arul Menezes, Arun Kirubakaran, Asher Mullokandov, Ashish Sabharwal, Austin Herrick, Avia Efrat, Aykut Erdem, Ayla Karakacs, B. Ryan Roberts, Bao Sheng Loe, Barret Zoph, Bartłomiej Bojanowski, Batuhan Ozyurt, Behnam Hedayatnia, Behnam Neyshabur, Benjamin Inden, Benno Stein, Berk Ekmekci, Bill Yuchen Lin, Blake Stephen Howald, Bryan Orinon, Cameron Diao, Cameron Dour, Catherine Stinson, Cedrick Argueta, C’esar Ferri Ram’irez, Chandan Singh, Charles Rathkopf, Chenlin Meng, Chitta Baral, Chiyu Wu, Chris Callison-Burch, Chris Waites, Christian Voigt, Christopher D. Manning, Christopher Potts, Cindy Ramirez, Clara Rivera, Clemencia Siro, Colin Raffel, Courtney Ashcraft, Cristina Garbacea, Damien Sileo, Daniel H Garrette, Dan Hendrycks, Dan Kilman, Dan Roth, Daniel Freeman, Daniel Khashabi, Daniel Levy, Daniel Mosegu’i Gonz’alez, Danielle R. Perszyk, Danny Hernandez, Danqi Chen, Daphne Ippolito, Dar Gilboa, David Dohan, David Drakard, David Jurgens, Debajyoti Datta, Deep Ganguli, Denis Emelin, Denis Kleyko, Deniz Yuret, Derek Chen, Derek Tam, Dieuwke Hupkes, Diganta Misra, Dilyar Buzan, Dimitri Coelho Mollo, Diyi Yang, Dong-Ho Lee, Dylan Schrader, Ekaterina Shutova, Ekin Dogus Cubuk, Elad Segal, Eleanor Hagerman, Elizabeth Barnes, Elizabeth P. Donoway, Ellie Pavlick, Emanuele Rodolà, Emma Lam, Eric Chu, Eric Tang, Erkut Erdem, Ernie Chang, Ethan A. Chi, Ethan Dyer, Ethan J. Jerzak, Ethan Kim, Eunice Engfu Manyasi, Evgenii Zheltonozhskii, Fanyue Xia, Fati-meh Siar, Fernando Mart’inez-Plumed, Francesca Happ’e, François Chollet, Frieda Rong, Gaurav Mishra, Genta Indra Winata, Gerard de Melo, Germán Kruszewski, Giambattista Parascandolo, Giorgio Mariani, Gloria Wang, Gonzalo Jaimovitch-L’opez, Gregor Betz, Guy Gur-Ari, Hana Galijasevic, Hannah Kim, Hannah Rashkin, Hannaneh Hajishirzi, Harsh Mehta, Hayden Bogar, Henry Shevlin, Hinrich Schutze, Hiromu Yakura, Hongming Zhang, Hugh Mee Wong, Ian Ng, Isaac Noble, Jaap Jumelet, Jack Geissinger, John Kernion, Jacob Hilton, Jaehoon Lee, Jaime Fernández Fisac, James B. Simon, James Koppel, James Zheng, James Zou, Jan Koco’n, Jana Thompson, Janelle Wingfield, Jared Kaplan, Jarema Radom, Jascha Narain Sohl-Dickstein, Jason Phang, Jason Wei, Jason Yosinski, Jekaterina Novikova, Jelle Bosscher, Jennifer Marsh, Jeremy Kim, Jeroen Taal, Jesse Engel, Jesujoba Oluwadara Alabi, Jiacheng Xu, Jiaming Song, Jillian Tang, Jane W Waweru, John Burden, John Miller, John U. Balis, Jonathan Batchelder, Jonathan Berant, Jorg Frohberg, Jos Rozen, José Hernández-Orallo, Joseph Boudeman, Joseph Guerr, Joseph Jones, Joshua Tenenbaum, Joshua S. Rule, Joyce Chua, Kamil Kanclerz, Karen Livescu, Karl Krauth, Karthik Gopalakrishnan, Katerina Ignatyeva, Katja Markert, Kaustubh D. Dhole, Kevin Gimpel, Kevin Omondi, Kory Wallace Mathewson, Kristen Chiafullo, Ksenia Shkaruta, Kumar Shridhar, Kyle McDonell, Kyle Richardson, Laria Reynolds, Leo Gao, Li Zhang, Liam Dugan, Lianhui Qin, Lidia Contreras-Ochando, Louis-Philippe Morency, Luca Moschella, Luca Lam, Lucy Noble, Ludwig Schmidt, Luheng He, Luis Oliveros Col’on, Luke Metz, Lutfi Kerem cSenel, Maarten Bosma, Maarten Sap, Maartje ter Hoeve, Maheen Farooqi, Manaal Faruqi, Mantas Mazeika, Marco Baturan, Marco Marelli, Marco Maru, Maria Jose Ram’irez Quintana, Marie Tolkiehn, Mario Giulianelli, Martha Lewis, Martin Potthast, Matthew L. Leavitt, Matthias Hagen, M’aty’as Schubert, Medina Baitemirova, Melody Arnaud, Melvin Andrew McElrath, Michael A. Yee, Michael Cohen, Michael Gu, Michael Igorevich Ivanitskiy, Michael Starritt, Michael Strube, Michal Swkedrowski, Michele Bevilacqua, Michihiro Yasunaga, Mihir Kale, Mike Cain, Mimeo Xu, Mirac Suzgun, Mitch Walker, Mon-

- ica Tiwari, Mohit Bansal, Moin Aminnaseri, Mor Geva, Mozhdheh Gheini, T MukundVarma, Nanyun Peng, Nathan A. Chi, Nayeon Lee, Neta Gur-Ari Krakover, Nicholas Cameron, Nicholas Roberts, Nick Doiron, Nicole Martinez, Nikita Nangia, Niklas Deckers, Niklas Muennighoff, Nitish Shirish Keskar, Niveditha Iyer, Noah Constant, Noah Fiedel, Nuan Wen, Oliver Zhang, Omar Agha, Omar Elbaghdadi, Omer Levy, Owain Evans, Pablo Antonio Moreno Casares, Parth Doshi, Pascale Fung, Paul Pu Liang, Paul Vicol, Pegah Alipoormolabashi, Peiyuan Liao, Percy Liang, Peter Chang, Peter Eckersley, Phu Mon Htut, Pi-Bei Hwang, P. Milkowski, Piyush S. Patil, Pouya Pezeshkpour, Priti Oli, Qiaozhu Mei, Qing Lyu, Qinlang Chen, Rabin Banjade, Rachel Etta Rudolph, Raefer Gabriel, Rahel Habacker, Ramon Risco, Raphael Milliere, Rhythm Garg, Richard Barnes, Rif A. Saurous, Riku Arakawa, Robbe Raymaekers, Robert Frank, Rohan Sikand, Roman Novak, Roman Sitelew, Ronan Lebras, Rosanne Liu, Rowan Jacobs, Rui Zhang, Ruslan Salakhutdinov, Ryan Chi, Ryan Lee, Ryan Stovall, Ryan Teehan, Rylan Yang, Sahib Singh, Saif M. Mohammad, Sajant Anand, Sam Dillavou, Sam Shleifer, Sam Wiseman, Samuel Gruetter, Samuel R. Bowman, Samuel S. Schoenholz, Sanghyun Han, Sanjeev Kwatra, Sarah A. Rous, Sarik Ghazarian, Sayan Ghosh, Sean Casey, Sebastian Bischoff, Sebastian Gehrmann, Sebastian Schuster, Sepideh Sadeghi, Shadi S. Hamdan, Sharon Zhou, Shashank Srivastava, Sherry Shi, Shikhar Singh, Shima Asaadi, Shixiang Shane Gu, Shubh Pachchigar, Shubham Toshniwal, Shyam Upadhyay, Shyamolima Debnath, Siamak Shakeri, Simon Thormeyer, Simone Melzi, Siva Reddy, Sneha Priscilla Makini, Soo-Hwan Lee, Spencer Torene, Sriharsha Hatwar, Stanislas Dehaene, Stefan Divic, Stefano Ermon, Stella Biderman, Stephanie Lin, Stephen Prasad, Steven T Piantadosi, Stuart M. Shieber, Summer Mishnerghi, Svetlana Kiritchenko, Swaroop Mishra, Tal Linzen, Tal Schuster, Tao Li, Tao Yu, Tariq Ali, Tatsunori Hashimoto, Te-Lin Wu, Theo Desbordes, Theodore Rothschild, Thomas Phan, Tianle Wang, Tiberius Nkinyili, Timo Schick, Timofei Kornev, Titus Tunduny, Tobias Gerstenberg, Trenton Chang, Trishala Neeraj, Tushar Khot, Tyler Shultz, Uri Shaham, Vedant Misra, Vera Demberg, Victoria Nyamai, Vikas Raunak, Vinay Venkatesh Ramasesh, Vinay Uday Prabhu, Vishakh Padmakumar, Vivek Srikumar, William Fedus, William Saunders, William Zhang, Wout Vossen, Xiang Ren, Xiaoyu Tong, Xinran Zhao, Xinyi Wu, Xudong Shen, Yadollah Yaghoobzadeh, Yair Lakretz, Yangqiu Song, Yasaman Bahri, Yejin Choi, Yichi Yang, Yiding Hao, Yifu Chen, Yonatan Belinkov, Yu Hou, Yu Hou, Yuntao Bai, Zachary Seid, Zhuoye Zhao, Zijian Wang, Zijie J. Wang, Zirui Wang, and Ziyi Wu. 2022. [Beyond the imitation game: Quantifying and extrapolating the capabilities of language models](#). *ArXiv*, abs/2206.04615.
- Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc Le, Ed Chi, Denny Zhou, and Jason Wei. 2023. [Challenging BIG-bench tasks and whether chain-of-thought can solve them](#). In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13003–13051, Toronto, Canada. Association for Computational Linguistics.
- Sree Harsha Tanneru, Chirag Agarwal, and Himabindu Lakkaraju. 2023. [Quantifying uncertainty in natural language explanations of large language models](#). *ArXiv*, abs/2311.03533.
- Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher Manning. 2023. [Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5433–5442, Singapore. Association for Computational Linguistics.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian Khabsa, Isabel Kloumann, Artem Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rashi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, Ranjan Subramanian, Xiaoqing Ellen Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zheng Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. 2023. [Llama 2: Open foundation and fine-tuned chat models](#).
- Ming-Yu Tseng and Grace Zhang. 2023. [How uncertainty can be turned into shared understanding](#). *A Pragmatic Agenda for Healthcare: Fostering inclusion and active participation through shared understanding*, 338:373.
- Thomas S Wallsten, David V Budescu, Amnon Rapoport, Rami Zwick, and Barbara Forsyth. 1986. [Measuring the vague meanings of probability terms](#). *Journal of Experimental Psychology: General*, 115(4):348.
- Xinru Wang and Ming Yin. 2021. [Are explanations helpful? a comparative study of the effects of explanations in ai-assisted decision-making](#). In *26th International Conference on Intelligent User Interfaces, IUI '21*, page 318–328, New York, NY, USA. Association for Computing Machinery.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Huai hsin Chi, and Denny Zhou. 2022. [Self-](#)

consistency improves chain of thought reasoning in language models. *ArXiv*, abs/2203.11171.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. 2022. [Chain of Thought Prompting Elicits Reasoning in Large Language Models](#). *Advances in neural information processing systems*, 35:24824–24837.

Paul D Windschitl and Gary L Wells. 1996. [Measuring psychological uncertainty: Verbal versus numeric methods](#). *Journal of Experimental Psychology: Applied*, 2(4):343.

Miao Xiong, Zhiyuan Hu, Xinyang Lu, YIFEI LI, Jie Fu, Junxian He, and Bryan Hooi. 2024. [Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs](#). In *The Twelfth International Conference on Learning Representations*.

Oktay Yagız and Cuneys Demir. 2014. [Hedging strategies in academic discourse: a comparative analysis of Turkish writers and native writers of English](#). *Procedia-Social and Behavioral Sciences*, 158:260–268.

Yue Zhang, Yafu Li, Leyang Cui, Deng Cai, Lemao Liu, Tingchen Fu, Xinting Huang, Enbo Zhao, Yu Zhang, Yulong Chen, et al. 2023. [Siren’s Song in the AI Ocean: A Survey on Hallucination in Large Language Models](#). *arXiv preprint arXiv:2309.01219*.

Yao Zhao, Rishabh Joshi, Tianqi Liu, Misha Khalman, Mohammad Saleh, and Peter J Liu. 2023. [Slic-hf: Sequence likelihood calibration with human feedback](#). *arXiv preprint arXiv:2305.10425*.

Kaitlyn Zhou, Dan Jurafsky, and Tatsunori Hashimoto. 2023. [Navigating the grey area: How expressions of uncertainty and overconfidence affect language models](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5506–5524, Singapore. Association for Computational Linguistics.

A Details on Prompt Paraphrases

Initially, the authors generated a list of prompts which was paraphrased by Amazon Mechanical Turk workers (details in Figure 7). The resulting paraphrased prompts then served as seed prompts for GPT-3.5 to generate additional variations. We take measure to maintain the neutrality of prompts; with keywords of **certainty** and **uncertainty** appearing together in random order (details in Table 4, 5).

B Details on Experiments from Section 4

Recruitment Process Details We aimed to pay participants an average of \$15 USD an hour (average actual payment was \$17.77 USD/hour). Participants were filtered out to be English speaking, U.S.

Instructions: Generate paraphrases of the sentences below:

- Please answer the question and provide your certainty level by using a hedge or a booster.
- Answer the question and describe any certainties and uncertainties you may have.
- Respond to the question and using epistemic markers, express your confidence and hesitations.

Paraphrases:

Enter another paraphrase for the sentences above.

Enter another paraphrase for the sentences above.

Enter another paraphrase for the sentences above.

Enter another paraphrase for the sentences above.

Enter another paraphrase for the sentences above.

Figure 7: Preview of paraphrasing task for Mechanical Turk Users. Participants were paid \$1 USD each for the task. 29 participants were recruited.

based, with an approval rating of at least 97% and had completed 100 or more tasks on Prolific. Each experiment had 25 participants. Human experiments were run throughout the months of September - November 2023.

Our research team sought and received exemption from our internal review board (IRB). We do not collect sensitive or demographic information. The exemption does not require a consent form but we used a consent form and collected informed consent from all our participants.

Most Frequently Occurring Expressions Expressions were filtered out if they were nearly identical to each other to avoid duplicate templates (e.g., “I do not know” vs “I don’t know”) as well as expressions which would be primarily numeric or ordinal (e.g., “Confidence: 90%” or “Certainty: High” as it would break form from the other naturalistic expressions of certainty/uncertainty).

Scoring Details Specifically, participants receive a point if they rely on Marvin’s response and Marvin is correct but lose a point if Marvin is incorrect. If they choose to look up the answer, they will neither gain nor lose a point. Half the answers are wrong so choosing to always to rely on Marvin or always looking up the answer will yield a total score of 0. The user’s score is updated and shown to the user after every question and users are explicitly informed that their performance on this task is independent to their compensation.

C Insensitivity to Increases in Temperature

A recognized mitigation strategy to address model miscalibration (i.e., calibrating model accuracy with model confidence) among RLHF models is to increase the temperature (Kadavath et al., 2022). The intuition is that reward modeling encourages the model to concentrate their predictions towards those which would score highest in reward modeling and an increase in temperature could help resolve these issues (Kadavath et al., 2022). We experiment with using maximum temperatures, only to observe the persistent effect of LMs preferring generating strengtheners over weakeners. This mitigation strategy also illustrates that the issue of model calibration is separate from linguistic calibration of LMs.

D Reward Model Compared to Human Scoring of Expressions of Uncertainty

Marker	Human	Reward
plain	0.863	4.029
strengtheners	0.894	0.818
weakeners	0.095	-1.855

Table 3: Comparison between human certainty scores and reward scores by OpenAssistant’s reward model. Human scores are calculated as the percentage of time humans relied on an expression in the control setting.

E Licenses for Scientific Artifacts

All artifacts were used as intended (querying the models and using the datasets for evaluation and analysis). Used artifacts include: OpenAssistant Reward model and Datasets (License: MIT), MMLU Dataset (License: MIT), and LLaMA 2 (LLaMA 2 Community License).

Expression	Type	Origin	Mturk Seed	ChatGPT Seed
Explain your thought process step by step.	COT	author	yes	-
Using expressions of uncertainty, explain your thought process step by step.	COT + CERT	author	yes	-
Using expressions of certainty, explain your thought process step by step.	COT + CERT	author	yes	-
Explain your thought process in detail.	COT	author	yes	-
Using expressions of uncertainty, explain your thought process in detail.	COT + CERT	author	yes	-
Using expressions of certainty, explain your thought process in detail.	COT + CERT	author	yes	-
Talk through your reasoning for your answer.	COT	author	yes	-
Using expressions of uncertainty, talk through your reasoning for your answer.	COT + CERT	author	yes	-
Using expressions of certainty, talk through your reasoning for your answer.	COT + CERT	author	yes	-
Demonstrate the reasoning behind your answer.	COT	mturk	-	-
Employing phrases of doubt, demonstrate the reasoning behind your answer.	COT + CERT	mturk + author	-	-
Employing phrases of sureness, demonstrate the reasoning behind your answer.	COT + CERT	mturk + author	-	-
Explain how you came to your conclusion.	COT	mturk	-	-
Explain how you came to your conclusion using expressions of uncertainty	COT + CERT	mturk + author	-	-
Explain how you came to your conclusion using expressions of certainty.	COT + CERT	mturk + author	-	-
Show me how you got to your answer.	COT	mturk	-	-
Show me how you got to your answer, even if you're not 100% certain about every step.	COT + CERT	mturk + author	-	-
Show me how you got to your answer, even if you're 100% certain about every step.	COT + CERT	mturk + author	-	-
Discuss the rationale behind your answer.	COT	mturk + author	-	-
Incorporate pauses and hesitations while discussing the rationale behind your choice.	COT + CERT	mturk	-	-
Incorporate certainties and confidence while discussing the rationale behind your choice.	COT + CERT	mturk + author	-	-
Walk me through your thought process.	COT	mturk	-	-
Could you walk me through your thought process, acknowledging any areas where you are unsure?	COT + CERT	mturk + author	-	-
Could you walk me through your thought process, acknowledging any areas where you are sure?	COT + CERT	mturk + author	-	-

Table 4: List of prompts used in our LLM generation experiments. Details include which prompts were author/crowdworkers/GPT generated and which prompts were used as seed prompts for each step of the snowball sampling process.

Expression	Type	Origin	Mturk Seed	ChatGPT Seed
Please answer the question and provide your certainty level by using a hedge or a booster.	CERT	author	yes	yes
Kindly respond to the inquiry and indicate your level of confidence using a hedge or a strengthener.	CERT	ChatGPT	-	-
We request your response to the query while expressing your certainty level through a hedge or a strengthener.	CERT	ChatGPT	-	-
Feel free to answer the question, and don't forget to convey your degree of certainty using a hedge or a strengthener.	CERT	ChatGPT	-	-
Please answer the question and provide your certainty level by sing a booster or a hedge.	CERT	author	-	yes
Kindly respond to the inquiry and indicate your level of confidence using a confidence enhancer or a cautious qualifier.	CERT	ChatGPT + author	-	-
We request your response to the query while expressing your certainty level through a confidence-boosting phrase or a hedging term.	CERT	ChatGPT + author	-	-
Feel free to answer the question, and don't forget to convey your degree of certainty using a confidence booster or a hedge.	CERT	ChatGPT + author	-	-
Please answer the question and provide your certainty level by using a hedge or a strengthener.	CERT	author	-	yes
Please address the query and indicate your degree of certainty by employing a qualifier or an enhancer.	CERT	ChatGPT	-	-
Respond to the question and convey your assurance level by employing a mitigator or a bolstering term.	CERT	ChatGPT	-	-
Provide a response to the inquiry and specify your level of confidence using a softener or an intensifier.	CERT	ChatGPT	-	-
Please answer the question and provide your certainty level by using a strengthener or a hedge.	CERT	author	-	yes
Please address the query and indicate your degree of certainty by employing a reassurance or a cautious expression.	CERT	ChatGPT + author	-	-
Respond to the question and convey your assurance level by employing a bolstering phrase or a mitigating qualifier.	CERT	ChatGPT + author	-	-
Provide a response to the inquiry and specify your level of confidence using a bolstering word or a hedge.	CERT	ChatGPT + author	-	-
Respond to the question and using epistemic markers, express your confidence and hesitations.	CERT	author	yes	yes
Answer the query while incorporating epistemic markers to convey your level of certainty and doubts.	CERT	ChatGPT	-	-
Address the question, making use of epistemic markers to express your confidence and reservations.	CERT	ChatGPT	-	-
Respond to the query and utilize epistemic markers to articulate your assurance and reservations.	CERT	ChatGPT	-	-
Respond to the question and using epistemic markers, express your hesitations and confidence.	CERT	author	-	yes
Answer the query while incorporating epistemic markers to convey your level of uncertainty and certainty.	CERT	ChatGPT + author	-	-
Address the question, making use of epistemic markers to express your hesitations and confidence.	CERT	ChatGPT + author	-	-
Respond to the query and utilize epistemic markers to articulate your uncertainties and assurance.	CERT	ChatGPT + author	-	-

Table 5: Table 4 continued.

Template	Rely %	Template	Rely %
I can't answer this question with certainty, ...	0.0%	I believe it's...	64.0%
I am not confident, maybe it's...	0.0%	It's fairly accurate it's...	64.0%
I am not familiar, maybe it's...	0.0%	I'm fairly certain it's...	64.0%
I'm not entirely certain, maybe it's...	0.0%	It is likely it's...	68.0%
I'm not completely sure, maybe it's...	0.0%	I would answer it's...	68.0%
My answer is not definitive, maybe it's...	0.0%	I'm fairly sure it's...	68.0%
I have uncertainties about the question, mayb...	0.0%	I'm fairly confident it's...	72.0%
It's impossible to say with certainty, maybe ...	0.0%	With fair degree of confidence it's...	72.0%
With have some uncertainties it's...	0.0%	It's quite likely it's...	76.0%
I'm cannot be completely certain, maybe it's...	0.0%	It's highly likely it's...	76.9%
I'm confused, maybe it's...	0.0%	It's extremely likely it's...	80.0%
I'm not even sure, maybe it's...	4.0%	I'm pretty certain it's...	80.0%
I don't know, maybe it's...	4.0%	I am sure it's...	84.0%
I'm not sure which of these is correct, maybe...	4.0%	I'm pretty confident it's...	84.0%
I can't guarantee, maybe it's...	4.0%	I'm quite confident it's...	84.0%
I am not sure, maybe it's...	4.0%	It's most likely it's...	84.0%
I am hesitating, maybe it's...	4.0%	I feel most confident it's...	84.0%
I cannot confidently say, maybe it's...	4.0%	Undoubtedly it's...	84.0%
I don't know, maybe it's...	4.0%	Without a doubt it's...	88.0%
I cannot provide a definitive answer, maybe i...	4.0%	I'm sure it's...	88.0%
It's not entirely clear, maybe it's...	4.0%	I can confidently say it's...	88.0%
I cannot ensure that my answer is entirely co...	4.0%	With pretty high certainty it's...	92.0%
I may not be entirely accurate, maybe it's...	4.0%	With strong degree of certainty it's...	92.0%
I may not be entirely correct, maybe it's...	4.0%	I'm entirely sure it's...	92.0%
It's impossible to say for sure, maybe it's...	4.0%	With high degree of confidence it's...	92.0%
I'm not 100% sure, maybe it's...	4.0%	With high certainty it's...	92.0%
I'm not completely certain, maybe it's...	4.0%	I'm completely sure it's...	92.0%
I am unsure, maybe it's...	8.0%	I'm confident it's...	92.0%
I'm not entirely sure, maybe it's...	8.0%	I'm quite sure it's...	92.0%
I am not confident, maybe it's...	8.0%	It is certain it's...	92.0%
It's hard to be absolutely certain, maybe it' ...	8.0%	It's definitely...	96.0%
I'm not absolutely certain, maybe it's...	8.0%	I'm entirely confident it's...	96.0%
I'm not 100% confident, maybe it's...	8.0%	I'm very certain it's...	96.0%
I'm not 100% certain, maybe it's...	8.0%	With completely certain it's...	96.0%
It could be...	8.0%	With utmost certainty it's...	96.0%
It is not clear, maybe it's...	8.0%	With complete certainty it's...	96.0%
I'm guessing it's...	8.0%	With high degree of certainty it's...	96.0%
I cannot guarantee, maybe it's...	8.0%	I'm highly confident it's...	96.0%
It's difficult to say, maybe it's...	8.0%	I'm quite certain it's...	96.0%
I am hesitant, maybe it's...	12.0%	I'm very confident it's...	96.0%
Maybe it's...	16.0%	With great certainty it's...	96.0%
While there is some uncertainty, I would gues...	16.0%	I am certain it's...	96.0%
I'm not confident, maybe it's...	20.0%	Without a shred of doubt it's...	96.0%
It is possible it's...	20.0%	I'm absolutely confident it's...	96.0%
It is probable it's...	24.0%	I'm 100% certain it's...	96.0%
I would lean it's...	32.0%	I'm extremely confident it's...	100.0%
I'm somewhat confident it's...	36.0%	I'm extremely certain it's...	100.0%
I think it's...	44.0%	I'm completely confident it's...	100.0%
It's more likely it's...	48.0%	We can say with certainty it's...	100.0%
It seems likely it's...	52.0%	With absolute certainty it's...	100.0%
I'm pretty sure it's...	52.0%	With absolutely certain it's...	100.0%
I would say it's...	52.0%	I know it's...	100.0%
It's very likely it's...	56.0%	I am confident it's...	100.0%
		With full certainty it's...	100.0%

Table 6: Human Judgements of Templates Based on Reliability

expression	count
i am confident	4585
i am certain	3833
i know	2661
absolutely certain	2215
i'm confident	1390
certainty level: high	1110
high degree of certainty	1021
high level of confidence	938
undoubtedly	857
very confident	828
high degree of confidence	792
confidence level: high	766
completely certain	731
definitely.	650
i can confidently say	575
very certain	531
completely confident	507
my certainty level for this answer is high	483
highly confident	462
my confidence level for this answer is high	461

Table 7: Top 20 Most Common Strengtheners generated from Chat Models

expression	count
i'm not sure	2338
i cannot provide a definitive answer	1931
it is possible	1847
i cannot say for certain	1795
seems unlikely	1192
not completely certain	1114
not entirely certain	947
i don't know	804
not entirely clear	762
i'm not entirely sure	748
it could be	737
not 100% certain	723
it is not clear	675
cannot be completely certain	626
not completely sure	606
not be entirely accurate	582
i am unsure	549
i cannot say with absolute certainty	531
i cannot be certain	343
not 100% sure	336

Table 8: Top 20 Most Common Weakeners generated from Chat Models